

Welcome To This Notebook 🤗

Project Name :- Total Caps Production Predictor

Problem Statement :-

The goal of this project is to train a Machine Learning model that accurately estimates the total production of caps based on Date & Time and Production per hour features. The model aims to contribute to optimizing production planning and resource allocation.

Problems to Solve:-

- 1) Outlier Detection and outlier treatment of 'Production Per Hour' feature
- 2) Train a Machine Learning model that accurately estimates the total production of caps based on Date & Time and Production per hour features.

Dataset:

The dataset includes the following features:

Date & Time: The timestamp indicating when the production data was recorded.

Production per Hour: The amount of production recorded per hour.

Target Variable:

Total Production: The total production accumulated during the recorded period.

Importing Libraries

In [309...

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
from sklearn.model_selection import GridSearchCV
```

Importing the Dataset

In [310...

```
df1 = pd.read_csv(r"C:\Users\91770\Downloads\cap-production-data.csv")
```

Settiing Display options to ensure feature name visibility

```
In [311... pd.set_option('display.max_columns',None)
```

Warning Suppression

```
In [312... import warnings
warnings.filterwarnings('ignore')
```

Creating copy of original dataset

```
In [313... df = df1.copy()
```

Eyeballing Dataset

```
In [314... #Checking first 5 rows of dataset

df.head()
```

```
Out[314]:
```

	Date & Time	Production per hour	Total Production
0	27/01/2024 17:11:44	77	77
1	27/01/2024 18:11:44	52	129
2	27/01/2024 19:11:44	62	191
3	27/01/2024 20:11:44	47	238
4	27/01/2024 21:11:44	35	273

```
In [315... # Checking shape of dataset

df.shape
```

```
Out[315]: (499, 3)
```

```
In [316... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 499 entries, 0 to 498
Data columns (total 3 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Date & Time            499 non-null    object
 1   Production per hour    499 non-null    int64
 2   Total Production      499 non-null    object
dtypes: int64(1), object(2)
memory usage: 11.8+ KB
```

```
In [317... # Checing columns

df.columns
```

```
Out[317]: Index(['Date & Time', 'Production per hour', 'Total Production'], dtype='object')
```

```
In [318... # Checking datatypes
```

```
df.dtypes
```

```
Out[318]: Date & Time          object
Production per hour      int64
Total Production         object
dtype: object
```

```
In [319... # Checking Duplicated values
```

```
df.duplicated().sum()
```

```
Out[319]: 0
```

Initial Observations About Raw Data

- 1. We have 499 rows and 3 columns
Independent Feature : 'Date & Time' and 'Production per hour'

Dependent Feature : 'Total Production'
- 1. Datatypes of 'Date & Time' and 'Production per hour' are 'object' we have to change it to 'datetime' and 'int64' respectively.
- 1. We have to create some useful features using 'Date & Time' feature so we can enrich our dataset.
- 1. In Total production there is comma in some 'Total Production' feature, e.g.:-22,685, we have replace comma.
- 1. There are no duplicate values in the dataset
- 1. Rename feature name 'Production per hour' to 'Production Per Hour'.

Data Cleaning

```
In [320... # Replacing comma with space in Total Production variable
```

```
df['Total Production'] = df['Total Production'].str.replace(',', '')
```

```
In [321... # Changing datatype of 'Total Production Column'
```

```
df['Total Production'] = df['Total Production'].astype('int64')
```

```
In [322... # Changing Datatype of 'Date & Time' feature
```

```
df['Date & Time'] = pd.to_datetime(df['Date & Time'], format='%d/%m/%Y %H:%M:%S')
```

```
In [323... # Checking datatypes
```

```
df.dtypes
```

```
Out[323]: Date & Time          datetime64[ns]
Production per hour      int64
Total Production         int64
dtype: object
```

In [324...

Renaming feature name 'Production per hour' to 'Production Per Hour'.

df.rename(columns={'Production per hour': 'Production Per Hour'}, inplace=True)

df.sample()

Out[325]:

	Date & Time	Production Per Hour	Total Production
294	2024-02-08 23:11:44	63	21419

Enriching our dataset using 'Date & Time' Variable

Extracting Year, Month and Days

In [326...

df['Year'] = df['Date & Time'].dt.year

df['Month'] = df['Date & Time'].dt.month

df['Day'] = df['Date & Time'].dt.day

Extracting Day of Week

In [327...

df['Day_of_Week'] = df['Date & Time'].dt.dayofweek

Extracting Hour, Minutes and Seconds

In [328...

df['Hour'] = df['Date & Time'].dt.hour
df['Minute'] = df['Date & Time'].dt.minute
df['Second'] = df['Date & Time'].dt.second

Checking first 5 rows just to confirm about all features created
df.head()

In [329...

Out[329]:

	Date & Time	Production Per Hour	Total Production	Year	Month	Day	Day_of_Week	Hour	Minute	Second
0	2024-01-27 17:11:44	77	77	2024	1	27	5	17	11	44
1	2024-01-27 18:11:44	52	129	2024	1	27	5	18	11	44
2	2024-01-27 19:11:44	62	191	2024	1	27	5	19	11	44
3	2024-01-27 20:11:44	47	238	2024	1	27	5	20	11	44
4	2024-01-27 21:11:44	35	273	2024	1	27	5	21	11	44

```
In [330...] # Checking datatypes of newly created features

df.dtypes
```

```
Out[330]: Date & Time          datetime64[ns]
Production Per Hour      int64
Total Production         int64
Year                     int64
Month                    int64
Day                      int64
Day_of_Week              int64
Hour                     int64
Minute                   int64
Second                   int64
dtype: object
```

As we have extracted required information from 'Date & Time' feature, we will now drop 'Date & Time' feature from dataset because of below reason

If the date and time features you extracted to capture the information you need for your analysis or modeling, keeping the original datetime column might introduce redundancy and unnecessary dimensionality.

```
In [331...] df = df.drop(['Date & Time'],axis=1)
```

Exploratory Data Analysis (EDA)

```
In [332...] df.describe()
```

```
Out[332]:
```

	Production Per Hour	Total Production	Year	Month	Day	Day_of_Week	Hour	Minu
count	499.000000	499.000000	499.0	499.000000	499.000000	499.000000	499.000000	499
mean	65.917836	18268.839679	2024.0	1.793587	13.004008	2.979960	11.474950	11
std	278.952637	8902.021523	0.0	0.405136	9.373550	2.001907	6.957767	C
min	-63.000000	77.000000	2024.0	1.000000	1.000000	0.000000	0.000000	11
25%	16.500000	11916.500000	2024.0	2.000000	6.000000	1.000000	5.000000	11
50%	46.000000	18970.000000	2024.0	2.000000	11.000000	3.000000	11.000000	11
75%	75.000000	25150.000000	2024.0	2.000000	16.000000	5.000000	18.000000	11
max	5879.000000	32893.000000	2024.0	2.000000	31.000000	6.000000	23.000000	11


Observation

There are outliers in 'Production Per Hour' feature.

```
In [333...] # Checking outliers for 'Production Per Hour' feature using boxplot

plt.figure(figsize=(20, 8))
```

```
sns.boxplot(x='Production Per Hour', data=df)
plt.show()
```



In [334...

```
# Creating heat map
```

```
selected_columns = ['Production Per Hour', 'Total Production', 'Month', 'Day', 'Day_of_Week', 'Hour', 'Minute', 'Second']

plt.figure(figsize=(20, 8))
sns.heatmap(df[selected_columns].corr(), annot=True, cmap='viridis')
plt.title('Heatmap for Selected Columns')
plt.show()
```



In [335...

```
# Checking Missing Values
```

```
df.isna().sum()
```

Out[335]:

```
Production Per Hour    0
Total Production      0
Year                  0
Month                 0
Day                  0
Day_of_Week           0
Hour                  0
Minute                0
Second                0
dtype: int64
```

Observation :

- 1. There are no any highly correlated feature which is greater than 0.9.

- 1. There are no any missing values in dataset

Feature Engineering

Train Test Split

```
In [336... X = df.drop(['Total Production'],axis=1)
y = df['Total Production']
```

```
In [337... X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.3, random_stat
```

```
In [338... X_train.shape
```

```
Out[338]: (349, 8)
```

```
In [339... X_test.shape
```

```
Out[339]: (150, 8)
```

Problem No. 1


To detect anamolies/outliers in 'Production Per Hour' feature and treat them.

Detecting of outlier :- Using Box Plot

Treatment for outlier :- Capping


```
In [340... # Checking outliers for 'Production Per Hour' feature using boxplot
```

```
plt.figure(figsize=(20,8))
sns.boxplot(x='Production Per Hour', data=df)
plt.show()
```



```
In [341... plt.figure(figsize=(16,5))
sns.distplot(df['Production Per Hour'])
```

```
Out[341]: <AxesSubplot:xlabel='Production Per Hour', ylabel='Density'>
```



Observation:-

As we can see 'Production Per Hour' follows skewed distribution, hence we will use upper limit and lower limit using IQR method to cap the outliers.

In [342... *# Calculating 25th and 75th percentile*

```
percentile25 = X_train['Production Per Hour'].quantile(0.25)
percentile75 = X_train['Production Per Hour'].quantile(0.75)
```

In [343... *# Calculating IQR*

```
IQR = percentile75 - percentile25
```

In [344... IQR

Out[344]: 61.0

In [345...
upper_limit = percentile75 + 1.5 * IQR
lower_limit = percentile25 - 1.5 * IQR

In [346...
print('Upper_limit',upper_limit)
print('Lower_limit',lower_limit)

```
Upper_limit 168.5
Lower_limit -75.5
```

Finding Outliers

In [347... X_train[X_train['Production Per Hour'] > upper_limit]

Out[347]:

	Production Per Hour	Year	Month	Day	Day_of_Week	Hour	Minute	Second
184	1532	2024	2	4	6	9	11	44
62	5879	2024	1	30	1	7	11	44
21	1002	2024	1	28	6	14	11	44
151	526	2024	2	3	5	0	11	44

In [348... X_train[X_train['Production Per Hour'] < lower_limit]

Out[348]:

	Production Per Hour	Year	Month	Day	Day_of_Week	Hour	Minute	Second
--	---------------------	------	-------	-----	-------------	------	--------	--------


Capping the anamolies/outliers

```
In [349... X_train['Production Per Hour'] = np.where(
    X_train['Production Per Hour'] > upper_limit,
    upper_limit,

    np.where(
        X_train['Production Per Hour'] < lower_limit,
        lower_limit,
        X_train['Production Per Hour']
    )
)
```

```
In [350... # Checking outliers for 'Production Per Hour' feature using boxplot after capping

plt.figure(figsize=(20,8))
sns.boxplot(x='Production Per Hour', data=X_train)
plt.show()
```



Problem No. 2 :

Model Training

```
In [351... def train_test_models(models, X_train, X_test, y_train, y_test):
    r2_scores = {}

    for model_name, model in models.items():
        model.fit(X_train, y_train)

        r2 = r2_score(y_test, model.predict(X_test))

        r2_scores[model_name] = r2

        print(f"{model_name} R2 Score: {r2}")

    return r2_scores
```

```
In [352... models = {
    'Decision Tree':DecisionTreeRegressor(),
```

```
'Random Forest': RandomForestRegressor()
}
```

In [353...

```
results = train_test_models(models, X_train, X_test, y_train, y_test)
```

Decision Tree R2 Score: 0.9998756544197035

Random Forest R2 Score: 0.99956006932605

Observation:-

We are getting 0.99 r2 score for both models

Hyperparameter Tuning

In [354...

```
def hyperparameter_tuning(models, param_grids, X_train, X_test, y_train, y_test):
    best_params = {}
    best_scores = {}

    for model_name, model in models.items():
        param_grid = param_grids[model_name]

        grid_search = GridSearchCV(model, param_grid, scoring='r2', cv=5)
        grid_search.fit(X_train, y_train)

        best_params[model_name] = grid_search.best_params_
        best_scores[model_name] = grid_search.best_score_

        best_estimator = grid_search.best_estimator_
        test_predictions = best_estimator.predict(X_test)
        test_score = r2_score(y_test, test_predictions)

        print(f"{model_name} - Best Parameters: {best_params[model_name]}")
        print(f"{model_name} - Best r2 Score: {best_scores[model_name]}")
        print(f"{model_name} - Test Score: {test_score}\n")

    return best_params, best_scores
```

In [355...

```
param_grids = {
    'Decision Tree': {
        'criterion': ['mse', 'mae'],
        'max_depth': [None, 5, 10, 15],
        'min_samples_split': [2, 5, 10]
    },
    'Random Forest': {
        'n_estimators': [50, 100, 150],
        'criterion': ['mse', 'mae'],
        'max_depth': [None, 5, 10, 15],
        'min_samples_split': [2, 5, 10]
    }
}
```

In [356...

```
best_params, best_scores = hyperparameter_tuning(models, param_grids, X_train, X_test, y_train, y_test)
```

Decision Tree - Best Parameters: {'criterion': 'mse', 'max_depth': 10, 'min_samples_split': 2}

Decision Tree - Best r2 Score: 0.9986768282159989

Decision Tree - Test Score: 0.9998566041718395

Random Forest - Best Parameters: {'criterion': 'mae', 'max_depth': None, 'min_samples_split': 2, 'n_estimators': 150}

Random Forest - Best r2 Score: 0.9987925711595385

Random Forest - Test Score: 0.9995921472740097

Observation :

Both models are giving good prediction

Thank You For Going Through This Notebook 😊

